

```

"""HTTP server base class.
1
2
Note: the class in this module doesn't implement any HTTP request; see
3
SimpleHTTPServer for simple implementations of GET, HEAD and POST
4
(including CGI scripts). It does, however, optionally implement HTTP/1.1
5
persistent connections, as of version 0.3.
6
7
Contents:
8
9
- BaseHTTPRequestHandler: HTTP request handler base class
10
- test: test function
11
12
XXX To do:
13
14
- log requests even later (to capture byte count)
15
- log user-agent header and other interesting goodies
16
- send error log to separate file
17
18
19
20
# See also:
21
#
22
# HTTP Working Group
23
# INTERNET-DRAFT
24
# <draft-ietf-http-v10-spec-00.txt>
25
# Expires September 8, 1995
26
#
27
# URL: http://www.ics.uci.edu/pub/ietf/http/draft-ietf-http-v10-spec-00.txt
28
#
29
# and
30
#
31
# Network Working Group
32
# Request for Comments: 2616
33
# Obsoletes: 2068
34
# Category: Standards Track
35
#
36
# URL: http://www.faqs.org/rfcs/rfc2616.html
37
38
# Log files
39
# -----
40
#
41
# Here's a quote from the NCSA httpd docs about log file format.
42
#
43
# | The logfile format is as follows. Each line consists of:
44
# |
45
# | host rfc931 authuser [DD/Mon/YYYY:hh:mm:ss] "request" ddd bbbb
46
# |
47
# | host: Either the DNS name or the IP number of the remote client
48
# | rfc931: Any information returned by identd for this person,
49
# | - otherwise.
50
# | authuser: If user sent a userid for authentication, the user
51
# | name,
52
# | - otherwise.
53
# | DD: Day
54
# | Mon: Month (calendar name)
55
# | YYYY: Year
56
# | hh: hour (24-hour format, the machine's timezone)
57
# | mm: minutes

```

```

58 # |         ss: seconds
59 # |         request: The first line of the HTTP request as sent by the
        client.
60 # |         ddd: the status code returned by the server, - if not available.
61 # |         bbbb: the total number of bytes sent,
62 # |             *not including the HTTP/1.0 header*, - if not available
63 # |
64 # | You can determine the name of the file accessed through request.
65 #
66 # (Actually, the latter is only true if you know the server configuration
67 # at the time the request was made!)
68
69 __version__ = "0.3"
70
71 __all__ = ["HTTPServer", "BaseHTTPRequestHandler"]
72
73 import sys
74 import time
75 import socket # For gethostbyaddr()
76 from warnings import filterwarnings, catch_warnings
77 with catch_warnings():
78     if sys.py3kwarning:
79         filterwarnings("ignore", ".*mimertools has been removed",
80                       DeprecationWarning)
81     import mimertools
82 import SocketServer
83
84 # Default error message template
85 DEFAULT_ERROR_MESSAGE = """\
86 <head>
87 <title>Error response</title>
88 </head>
89 <body>
90 <h1>Error response</h1>
91 <p>Error code %(code)d.
92 <p>Message: %(message)s.
93 <p>Error code explanation: %(code)s = %(explain)s.
94 </body>
95 """
96
97 DEFAULT_ERROR_CONTENT_TYPE = "text/html"
98
99 def _quote_html(html):
100     return html.replace("&", "&amp;").replace("<", "&lt;").replace(">", "&gt;")
101
102 class HTTPServer(SocketServer.TCPServer):
103
104     allow_reuse_address = 1    # Seems to make sense in testing environment
105
106     def server_bind(self):
107         """Override server_bind to store the server name."""
108         SocketServer.TCPServer.server_bind(self)
109         host, port = self.socket.getsockname()[:2]
110         self.server_name = socket.getfqdn(host)
111         self.server_port = port
112
113

```

```

class BaseHTTPRequestHandler(SocketServer.StreamRequestHandler):
    """HTTP request handler base class.

    The following explanation of HTTP serves to guide you through the
    code as well as to expose any misunderstandings I may have about
    HTTP (so you don't need to read the code to figure out I'm wrong
    :-).

    HTTP (HyperText Transfer Protocol) is an extensible protocol on
    top of a reliable stream transport (e.g. TCP/IP). The protocol
    recognizes three parts to a request:

    1. One line identifying the request type and path
    2. An optional set of RFC-822-style headers
    3. An optional data part

    The headers and data are separated by a blank line.

    The first line of the request has the form

    <command> <path> <version>

    where <command> is a (case-sensitive) keyword such as GET or POST,
    <path> is a string containing path information for the request,
    and <version> should be the string "HTTP/1.0" or "HTTP/1.1".
    <path> is encoded using the URL encoding scheme (using %xx to signify
    the ASCII character with hex code xx).

    The specification specifies that lines are separated by CRLF but
    for compatibility with the widest range of clients recommends
    servers also handle LF. Similarly, whitespace in the request line
    is treated sensibly (allowing multiple spaces between components
    and allowing trailing whitespace).

    Similarly, for output, lines ought to be separated by CRLF pairs
    but most clients grok LF characters just fine.

    If the first line of the request has the form

    <command> <path>

    (i.e. <version> is left out) then this is assumed to be an HTTP
    0.9 request; this form has no optional headers and data part and
    the reply consists of just the data.

    The reply form of the HTTP 1.x protocol again has three parts:

    1. One line giving the response code
    2. An optional set of RFC-822-style headers
    3. The data

    Again, the headers and data are separated by a blank line.

    The response code line has the form

    <version> <responsecode> <responsestring>

```

```

172     where <version> is the protocol version ("HTTP/1.0" or "HTTP/1.1"),
173     <responsecode> is a 3-digit response code indicating success or
174     failure of the request, and <responsestring> is an optional
175     human-readable string explaining what the response code means.
176
177     This server parses the request and the headers, and then calls a
178     function specific to the request type (<command>). Specifically,
179     a request SPAM will be handled by a method do_SPAM(). If no
180     such method exists the server sends an error response to the
181     client. If it exists, it is called with no arguments:
182
183     do_SPAM()
184
185     Note that the request name is case sensitive (i.e. SPAM and spam
186     are different requests).
187
188     The various request details are stored in instance variables:
189
190     - client_address is the client IP address in the form (host,
191       port);
192
193     - command, path and version are the broken-down request line;
194
195     - headers is an instance of mimetools.Message (or a derived
196       class) containing the header information;
197
198     - rfile is a file object open for reading positioned at the
199       start of the optional input data part;
200
201     - wfile is a file object open for writing.
202
203     IT IS IMPORTANT TO ADHERE TO THE PROTOCOL FOR WRITING!
204
205     The first thing to be written must be the response line. Then
206     follow 0 or more header lines, then a blank line, and then the
207     actual data (if any). The meaning of the header lines depends on
208     the command executed by the server; in most cases, when data is
209     returned, there should be at least one header line of the form
210
211     Content-type: <type>/<subtype>
212
213     where <type> and <subtype> should be registered MIME types,
214     e.g. "text/html" or "text/plain".
215
216     """
217
218     # The Python system version, truncated to its first component.
219     sys_version = "Python/" + sys.version.split()[0]
220
221     # The server software version. You may want to override this.
222     # The format is multiple whitespace-separated strings,
223     # where each string is of the form name[/version].
224     server_version = "BaseHTTP/" + __version__
225
226     # The default request version. This only affects responses up until
227     # the point where the request line is parsed, so it mainly decides what
228     # the client gets back when sending a malformed request line.
229     # Most web servers default to HTTP 0.9, i.e. don't send a status line.

```

```

default_request_version = "HTTP/0.9"
230
231
def parse_request(self):
232
    """Parse a request (internal).
233
234
    The request should be stored in self.raw_requestline; the results
235
    are in self.command, self.path, self.request_version and
236
    self.headers.
237
238
    Return True for success, False for failure; on failure, an
239
    error is sent back.
240
241
    """
242
    self.command = None # set in case of error on the first line
243
    self.request_version = version = self.default_request_version
244
    self.close_connection = 1
245
    requestline = self.raw_requestline
246
    requestline = requestline.rstrip('\r\n')
247
    self.requestline = requestline
248
    words = requestline.split()
249
    if len(words) == 3:
250
        command, path, version = words
251
        if version[:5] != 'HTTP/':
252
            self.send_error(400, "Bad request version (%r)" % version)
253
            return False
254
        try:
255
            base_version_number = version.split('/', 1)[1]
256
            version_number = base_version_number.split(".")
257
            # RFC 2145 section 3.1 says there can be only one "." and
258
            # - major and minor numbers MUST be treated as
259
            # - separate integers;
260
            # - HTTP/2.4 is a lower version than HTTP/2.13, which in
261
            # - turn is lower than HTTP/12.3;
262
            # - Leading zeros MUST be ignored by recipients.
263
            if len(version_number) != 2:
264
                raise ValueError
265
            version_number = int(version_number[0]), int(version_number
266
                [1])
267
        except (ValueError, IndexError):
268
            self.send_error(400, "Bad request version (%r)" % version)
269
            return False
270
        if version_number >= (1, 1) and self.protocol_version >= "HTTP
            /1.1":
271
            self.close_connection = 0
272
        if version_number >= (2, 0):
273
            self.send_error(505,
274
                "Invalid HTTP Version (%s)" % base_version_number
            )
275
            return False
276
    elif len(words) == 2:
277
        command, path = words
278
        self.close_connection = 1
279
        if command != 'GET':
280
            self.send_error(400,
281
                "Bad HTTP/0.9 request type (%r)" % command)
282
            return False
283
    elif not words:
284
        return False

```

```

285         else:
286             self.send_error(400, "Bad request syntax (%r)" % requestline)
287             return False
288         self.command, self.path, self.request_version = command, path,
            version
289
290         # Examine the headers and look for a Connection directive
291         self.headers = self.MessageClass(self.rfile, 0)
292
293         conntype = self.headers.get('Connection', "")
294         if conntype.lower() == 'close':
295             self.close_connection = 1
296         elif (conntype.lower() == 'keep-alive' and
297               self.protocol_version >= "HTTP/1.1"):
298             self.close_connection = 0
299         return True
300
301     def handle_one_request(self):
302         """Handle a single HTTP request.
303
304         You normally don't need to override this method; see the class
305         __doc__ string for information on how to handle specific HTTP
306         commands such as GET and POST.
307
308         """
309         try:
310             self.raw_requestline = self.rfile.readline(65537)
311             if len(self.raw_requestline) > 65536:
312                 self.requestline = ''
313                 self.request_version = ''
314                 self.command = ''
315                 self.send_error(414)
316                 return
317             if not self.raw_requestline:
318                 self.close_connection = 1
319                 return
320             if not self.parse_request():
321                 # An error code has been sent, just exit
322                 return
323             mname = 'do_' + self.command
324             if not hasattr(self, mname):
325                 self.send_error(501, "Unsupported method (%r)" % self.
                    command)
326                 return
327             method = getattr(self, mname)
328             method()
329             self.wfile.flush() #actually send the response if not already
                done.
330         except socket.timeout, e:
331             #a read or a write timed out. Discard this connection
332             self.log_error("Request timed out: %r", e)
333             self.close_connection = 1
334             return
335
336     def handle(self):
337         """Handle multiple requests if necessary."""
338         self.close_connection = 1
339

```

```

        self.handle_one_request()
        while not self.close_connection:
            self.handle_one_request()

def send_error(self, code, message=None):
    """Send and log an error reply.

    Arguments are the error code, and a detailed message.
    The detailed message defaults to the short entry matching the
    response code.

    This sends an error response (so it must be called before any
    output has been generated), logs the error, and finally sends
    a piece of HTML explaining the error to the user.

    """

    try:
        short, long = self.responses[code]
    except KeyError:
        short, long = '???', '???'
    if message is None:
        message = short
    explain = long
    self.log_error("code %d, message %s", code, message)
    # using _quote_html to prevent Cross Site Scripting attacks (see
    # bug #1100201)
    content = (self.error_message_format %
               {'code': code, 'message': _quote_html(message), 'explain':
                explain})
    self.send_response(code, message)
    self.send_header("Content-Type", self.error_content_type)
    self.send_header('Connection', 'close')
    self.end_headers()
    if self.command != 'HEAD' and code >= 200 and code not in (204,
    304):
        self.wfile.write(content)

error_message_format = DEFAULT_ERROR_MESSAGE
error_content_type = DEFAULT_ERROR_CONTENT_TYPE

def send_response(self, code, message=None):
    """Send the response header and log the response code.

    Also send two standard headers with the server software
    version and the current date.

    """
    self.log_request(code)
    if message is None:
        if code in self.responses:
            message = self.responses[code][0]
        else:
            message = ''
    if self.request_version != 'HTTP/0.9':
        self.wfile.write("%s %d %s\r\n" %
                         (self.protocol_version, code, message))
    # print (self.protocol_version, code, message)

```

```

395         self.send_header('Server', self.version_string())
396         self.send_header('Date', self.date_time_string())
397
398     def send_header(self, keyword, value):
399         """Send a MIME header."""
400         if self.request_version != 'HTTP/0.9':
401             self.wfile.write("%s: %s\r\n" % (keyword, value))
402
403         if keyword.lower() == 'connection':
404             if value.lower() == 'close':
405                 self.close_connection = 1
406             elif value.lower() == 'keep-alive':
407                 self.close_connection = 0
408
409     def end_headers(self):
410         """Send the blank line ending the MIME headers."""
411         if self.request_version != 'HTTP/0.9':
412             self.wfile.write("\r\n")
413
414     def log_request(self, code='-', size='-'):
415         """Log an accepted request.
416
417         This is called by send_response().
418
419         """
420
421         self.log_message("%s" %s %s',
422                         self.requestline, str(code), str(size))
423
424     def log_error(self, format, *args):
425         """Log an error.
426
427         This is called when a request cannot be fulfilled. By
428         default it passes the message on to log_message().
429
430         Arguments are the same as for log_message().
431
432         XXX This should go to the separate error log.
433
434         """
435
436         self.log_message(format, *args)
437
438     def log_message(self, format, *args):
439         """Log an arbitrary message.
440
441         This is used by all other logging functions. Override
442         it if you have specific logging wishes.
443
444         The first argument, FORMAT, is a format string for the
445         message to be logged. If the format string contains
446         any % escapes requiring parameters, they should be
447         specified as subsequent arguments (it's just like
448         printf!).
449
450         The client ip address and current date/time are prefixed to every
451         message.
452

```



```

"""
453
454
sys.stderr.write("%s -- [%s] %s\n" %
455
                    (self.client_address[0],
456
                     self.log_date_time_string(),
457
                     format%args))
458
459
def version_string(self):
460
    """Return the server software version string."""
461
    return self.server_version + ' ' + self.sys_version
462
463
def date_time_string(self, timestamp=None):
464
    """Return the current date and time formatted for a message header
465
        """
466
    if timestamp is None:
467
        timestamp = time.time()
468
    year, month, day, hh, mm, ss, wd, y, z = time.gmtime(timestamp)
469
    s = "%s, %02d %3s %4d %02d:%02d:%02d GMT" % (
470
        self.weekdayname[wd],
471
        day, self.monthname[month], year,
472
        hh, mm, ss)
473
    return s
474
475
def log_date_time_string(self):
476
    """Return the current time formatted for logging."""
477
    now = time.time()
478
    year, month, day, hh, mm, ss, x, y, z = time.localtime(now)
479
    s = "%02d/%03s/%04d %02d:%02d:%02d" % (
480
        day, self.monthname[month], year, hh, mm, ss)
481
    return s
482
483
weekdayname = ['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun']
484
485
monthname = [None,
486
              'Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun',
487
              'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']
488
489
def address_string(self):
490
    """Return the client address formatted for logging.
491
492
    This version looks up the full hostname using gethostbyaddr(),
493
    and tries to find a name that contains at least one dot.
494
495
    """
496
497
    host, port = self.client_address[:2]
498
    return socket.getfqdn(host)
499
500
# Essentially static class variables
501
502
# The version of the HTTP protocol we support.
503
# Set this to HTTP/1.1 to enable automatic keepalive
504
protocol_version = "HTTP/1.0"
505
506
# The Message-like class used to parse headers
507
MessageClass = mimetools.Message
508
509
# Table mapping response codes to messages; entries have the

```

```

510 # form {code: (shortmessage, longmessage)}.
511 # See RFC 2616.
512 responses = {
513     100: ('Continue', 'Request received, please continue'),
514     101: ('Switching Protocols',
515         'Switching to new protocol; obey Upgrade header'),
516
517     200: ('OK', 'Request fulfilled, document follows'),
518     201: ('Created', 'Document created, URL follows'),
519     202: ('Accepted',
520         'Request accepted, processing continues off-line'),
521     203: ('Non-Authoritative Information', 'Request fulfilled from
522         cache'),
523     204: ('No Content', 'Request fulfilled, nothing follows'),
524     205: ('Reset Content', 'Clear input form for further input.'),
525     206: ('Partial Content', 'Partial content follows.'),
526
527     300: ('Multiple Choices',
528         'Object has several resources — see URI list'),
529     301: ('Moved Permanently', 'Object moved permanently — see URI
530         list'),
531     302: ('Found', 'Object moved temporarily — see URI list'),
532     303: ('See Other', 'Object moved — see Method and URL list'),
533     304: ('Not Modified',
534         'Document has not changed since given time'),
535     305: ('Use Proxy',
536         'You must use proxy specified in Location to access this '
537         'resource.'),
538     307: ('Temporary Redirect',
539         'Object moved temporarily — see URI list'),
540
541     400: ('Bad Request',
542         'Bad request syntax or unsupported method'),
543     401: ('Unauthorized',
544         'No permission — see authorization schemes'),
545     402: ('Payment Required',
546         'No payment — see charging schemes'),
547     403: ('Forbidden',
548         'Request forbidden — authorization will not help'),
549     404: ('Not Found', 'Nothing matches the given URI'),
550     405: ('Method Not Allowed',
551         'Specified method is invalid for this resource.'),
552     406: ('Not Acceptable', 'URI not available in preferred format.'),
553     407: ('Proxy Authentication Required', 'You must authenticate with
554         this proxy before proceeding.'),
555     408: ('Request Timeout', 'Request timed out; try again later.'),
556     409: ('Conflict', 'Request conflict.'),
557     410: ('Gone',
558         'URI no longer exists and has been permanently removed.'),
559     411: ('Length Required', 'Client must specify Content-Length.'),
560     412: ('Precondition Failed', 'Precondition in headers is false.'),
561     413: ('Request Entity Too Large', 'Entity is too large.'),
562     414: ('Request-URI Too Long', 'URI is too long.'),
563     415: ('Unsupported Media Type', 'Entity body in unsupported format
564         .'),
565     416: ('Requested Range Not Satisfiable',
566         'Cannot satisfy request range.'),

```

```

417: ('Expectation Failed',
    'Expect condition could not be satisfied.'),
500: ('Internal Server Error', 'Server got itself in trouble'),
501: ('Not Implemented',
    'Server does not support this operation'),
502: ('Bad Gateway', 'Invalid responses from another server/proxy
    .'),
503: ('Service Unavailable',
    'The server cannot process the request due to a high load'),
504: ('Gateway Timeout',
    'The gateway server did not receive a timely response'),
505: ('HTTP Version Not Supported', 'Cannot fulfill request.'),
}

def test(HandlerClass = BaseHTTPRequestHandler,
        ServerClass = HTTPServer, protocol="HTTP/1.0"):
    """Test the HTTP request handler class.

    This runs an HTTP server on port 8000 (or the first command line
    argument).

    """

    if sys.argv[1:]:
        port = int(sys.argv[1])
    else:
        port = 8000
    server_address = ('', port)

    HandlerClass.protocol_version = protocol
    httpd = ServerClass(server_address, HandlerClass)

    sa = httpd.socket.getsockname()
    print "Serving HTTP on", sa[0], "port", sa[1], "..."
    httpd.serve_forever()

if __name__ == '__main__':
    test()

```